

Case Study Brief

Hello! Thank you for applying with us as a backend developer. This mini project **should be completed within 5 days after you have received this document**. Please spare your time to complete this project with the best results. We are really pleased to answer your questions if there are unclear things.

Objective

Your mission is to build a backend service that automates the initial screening of a job application. The service will receive a candidate's CV and a project report, evaluate them against a specific job description and a case study brief, and produce a structured, AI-generated evaluation report.

Core Logic & Data Flow

The system operates with a clear separation of inputs and reference documents:

Candidate-Provided Inputs (The Data to be Evaluated):

1. **Candidate CV:** The candidate's resume (PDF).
2. **Project Report:** The candidate's project report to our take-home case study (PDF)

System-Internal Documents (The "Ground Truth" for Comparison):

1. **Job Description:** A document detailing the requirements and responsibilities for the role — You can use the job description you're currently applying. This document will be used as ground truth for **Candidate CV**.
 - To make sure the vector retrieval is accurate enough, you might need to ingest a few job description documents as well.
2. **Case Study Brief:** This document. Used as ground truth for **Project Report**. (PDF)
3. **Scoring Rubric:** A predefined set of parameters for evaluating CV and Report, each has its own documents. (PDF)

We want to see your ability to combine **backend engineering** with **AI workflows** (prompt design, LLM chaining, retrieval, resilience).

Deliverables

1. Backend Service (API endpoints)

Implement a backend service with at least the following RESTful API endpoints:

- `POST /upload`
 - Accepts multipart/form-data containing the Candidate CV and Project Report (PDF).
 - Stores these files, return each with its own ID for later processing.
- `POST /evaluate`
 - Triggers the asynchronous AI evaluation pipeline. Receives input job title (string), and **both** document ID.
 - Immediately returns a job ID to track the evaluation process.

```
{  
  "id": "456",  
  "status": "queued"  
}
```

- `GET /result/{id}`
 - Retrieves the status and result of an evaluation job. This endpoint should reflect the asynchronous, **multi-stage** nature of the process.
 - Possible responses:
 - While queued or processing

```
{  
  "id": "456",  
  "status": "queued" | "processing"  
}
```

- Once completed

```
{
  "id": "456",
  "status": "completed",
  "result": {
    "cv_match_rate": 0.82,
    "cv_feedback": "Strong in backend and cloud, limited AI integration experience...",
    "project_score": 4.5,
    "project_feedback": "Meets prompt chaining requirements, lacks error handling robustness...",
    "overall_summary": "Good candidate fit, would benefit from deeper RAG knowledge..."
  }
}
```

2. Evaluation Pipeline

Design and implement an AI-driven pipeline which will be triggered by **[POST] /evaluate** endpoint. Should consist these key Components:

- **RAG (Context Retrieval)**

- Ingest all **System-Internal Documents** (Job Description, Case Study Brief, Both Scoring Rubrics) into a vector database.
- Retrieve relevant sections and inject into prompts (e.g., “for CV scoring” vs “for project scoring”).

- **Prompt Design & LLM Chaining**

The pipeline should consists of

- CV Evaluation
 - Parse the candidate’s CV into structured data.
 - Retrieve relevant information from both Job Description and CV Scoring Rubrics.
 - Use an LLM to get these result: `cv_match_rate` & `cv_feedback`
- Project Report Evaluation
 - Parse the candidate’s Project Report into structured data.
 - Retrieve relevant information from both Case Study Brief and CV Scoring Rubrics.
 - Use an LLM to get these result: `project_score` & `project_feedback`
- Final Analysis
 - Use a final LLM call to synthesize the outputs from previous steps into a concise `overall_summary`.

- **Long-Running Process Handling**

- `POST /evaluate` should **not block** until LLM Chaining finishes.
- Store task, return job ID, allow `GET /result/{id}` to check later periodically.

- **Error Handling & Randomness Control**

- Simulate any edge cases you can think of and how well your service can handle them.
- Simulate failures from LLM API (timeouts, rate limit).
- Implement retries/back-off.
- Control LLM temperature or add validation layer to keep responses stable.

3. Standardized Evaluation Parameters

Define at least these scoring parameters:

CV Evaluation (Match Rate)

- **Technical Skills Match** (backend, databases, APIs, cloud, AI/LLM exposure).
- **Experience Level** (years, project complexity).
- **Relevant Achievements** (impact, scale).
- **Cultural Fit** (communication, learning attitude).

Project Deliverable Evaluation

- **Correctness** (meets requirements: prompt design, chaining, RAG, handling errors).
- **Code Quality** (clean, modular, testable).
- **Resilience** (handles failures, retries).

- **Documentation** (clear README, explanation of trade-offs).
- **Creativity / Bonus** (optional improvements like authentication, deployment, dashboards).

Each parameter can be scored **1–5**, then aggregated to final score.

Requirements

- Use any backend framework (Rails, Django, Node.js, etc.).
- Use a proper LLM service (e.g., OpenAI, Gemini, or OpenRouter). There are several free LLM API providers available.
- Use a simple **vector DB** (e.g. ChromaDB, Qdrant, etc) or **RAG-as-a-service** (e.g. Ragie, S3 Vector, etc), any of your own choice.
- Provide README with run instructions + explanation of design choices.
- Provide the documents together with their ingestion scripts in the repository for reproducibility purposes.